

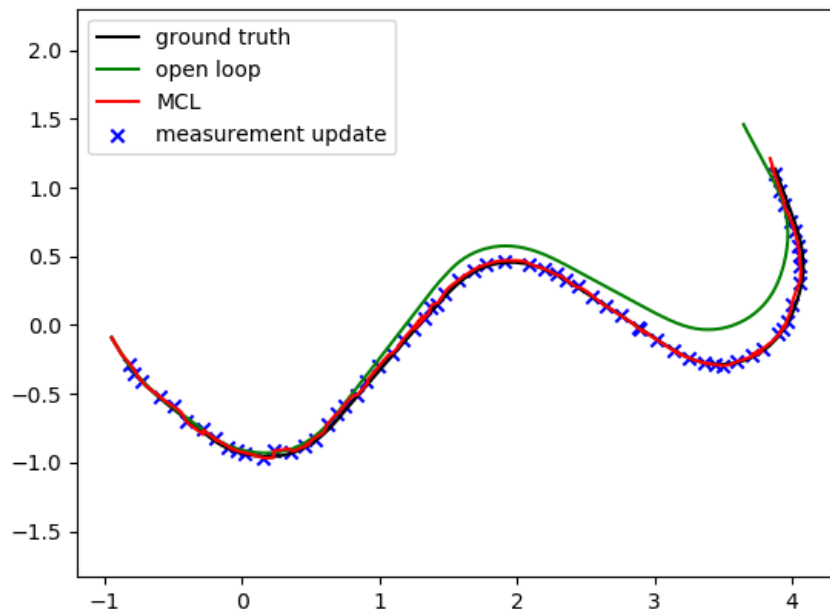
AA 274A: Principles of Robot Autonomy I

Problem Set 4

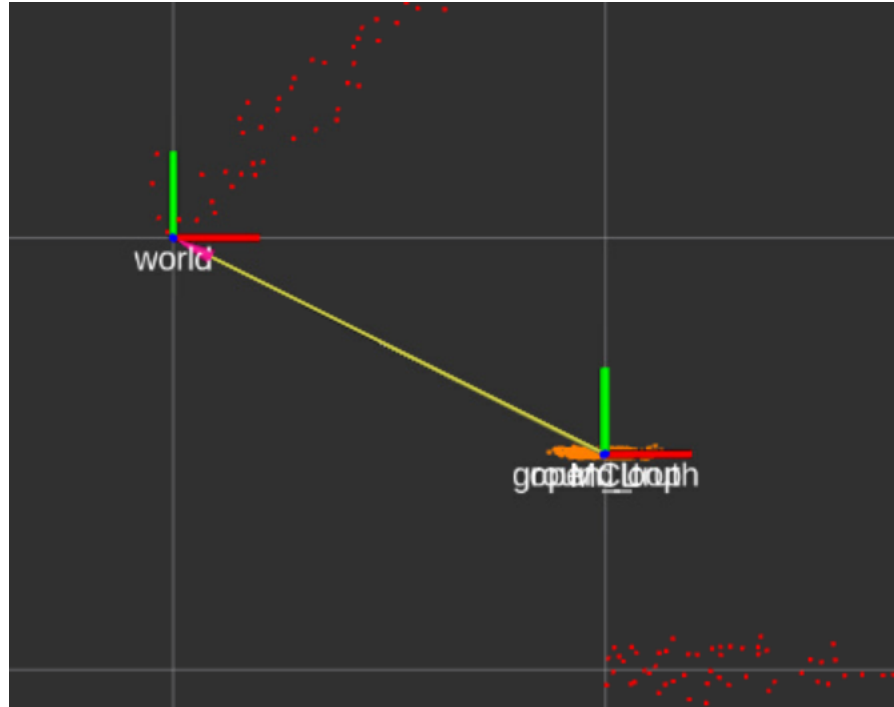
Name: Abhyudit Singh Manhas
SUID: 06645995

Extra Credit Problem: Monte Carlo Localization

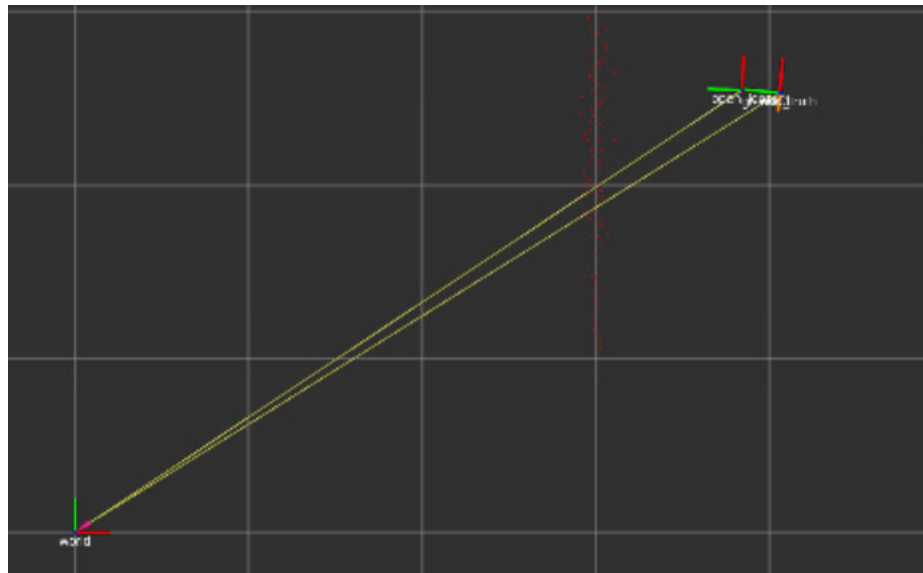
- (i) Implemented the transition update step of MCL by completing the `transition_model()` method in the `MonteCarloLocalization` class and the `transition_update()` method in the `ParticleFilter` class in `particle_filter.py`.
Validated my work by running `validate_transition_model()` from `validate_particle_filter.py`.
- (ii) Implemented the measurement update step by completing the `measurement_update()`, `measurement_model()`, `compute_innovations()`, and `compute_predicted_measurements()` methods of the `MonteCarloLocalization` class in `particle_filter.py`.
Validated my work by running `validate_predicted_measurements()` and `validate_compute_innovations()` from `validate_particle_filter.py`.
- (iii) Implemented the low variance sampling algorithm in the `resample()` method of the `ParticleFilter` class in `particle_filter.py`.
Validated my work by running `validate_resample()` and `validate_mc_localization()` from `validate_particle_filter.py`.
The file `mc_localization.png` is shown below.



- (iv) Just like in EKF, we observe that the MCL state estimates and ground truth diverge with sharp turns or collisions with the walls. But it eventually catches up by correcting for the error. With more particles, the MCL becomes more robust to sharp turns and wall collisions, i.e., the performance improves. The screenshots in RViz (for 1000 particles) are shown below:
- (1) The initial state:

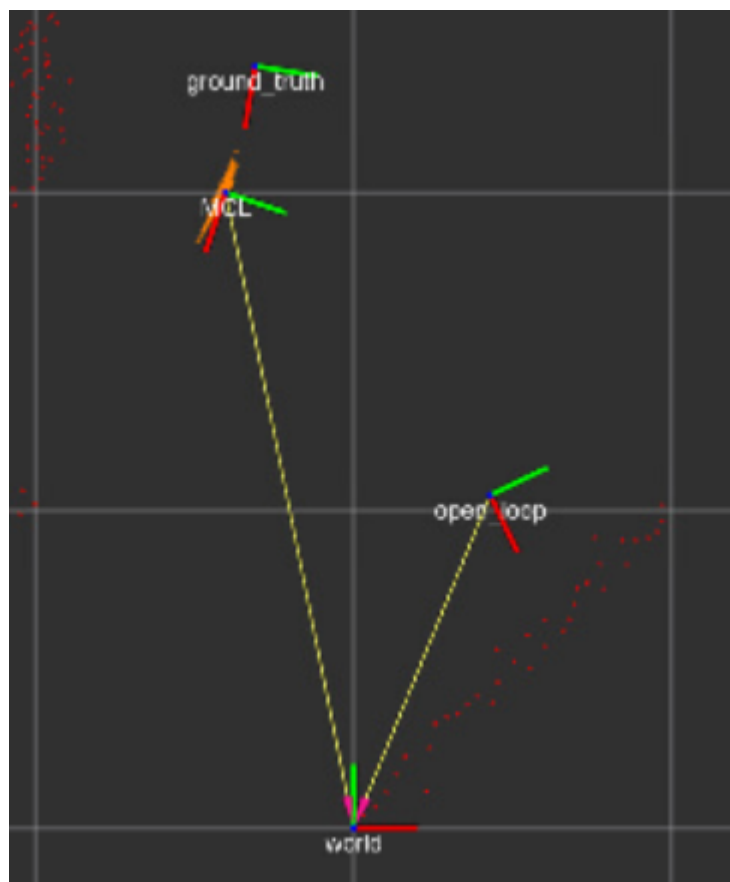


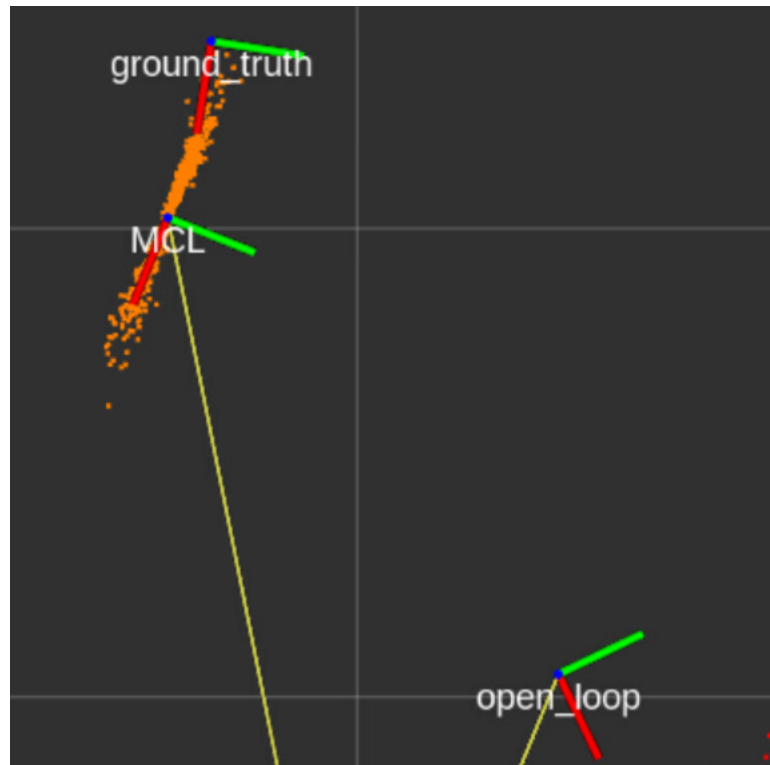
- (2) The TurtleBot has moved far from the initial state:





(3) The state estimates have diverged:





- (v) Vectorized all the functions and eliminated all loops. The entire program `particle_filter.py` is shown below:

```

1 import numpy as np
2 import scipy.linalg # You may find scipy.linalg.block_diag useful
3 import scipy.stats # You may find scipy.stats.multivariate_normal.pdf useful
4 from . import turtlebot_model as tb
5
6 EPSILON_OMEGA = 1e-3
7
8 class ParticleFilter(object):
9     """
10     Base class for Monte Carlo localization and FastSLAM.
11
12     Usage:
13     pf = ParticleFilter(x0, R)
14     while True:
15         pf.transition_update(u, dt)
16         pf.measurement_update(z, Q)
17         localized_state = pf.x
18     """
19
20     def __init__(self, x0, R):
21         """
22         ParticleFilter constructor.
23
24         Inputs:
25         x0: np.array[M,3] - initial particle states.
26         R: np.array[2,2] - control noise covariance (corresponding to dt = 1 second).
27         """
28         self.M = x0.shape[0] # Number of particles
29         self.xs = x0 # Particle set [M x 3]
30         self.ws = np.repeat(1. / self.M, self.M) # Particle weights (initialize to uniform) [M]
31         self.R = R # Control noise covariance (corresponding to dt = 1 second) [2 x 2]
32
33     @property
34     def x(self):
35         """
36         Returns the particle with the maximum weight for visualization.
37
38         Output:
39         x: np.array[3,] - particle with the maximum weight.
40         """
41         idx = self.ws == self.ws.max()
42         x = np.zeros(self.xs.shape[1:])
43         x[:2] = self.xs[idx, :2].mean(axis=0)
44         th = self.xs[idx, 2]
45         x[2] = np.arctan2(np.sin(th).mean(), np.cos(th).mean())
46         return x
47
48     def transition_update(self, u, dt):
49         """
50         Performs the transition update step by updating self.xs.
51
52         Inputs:
53         u: np.array[2,] - zero-order hold control input.
54         dt: float - duration of discrete time step.
55
56         Output:
57         None - internal belief state (self.xs) should be updated.
58         """
59         ##### Code starts here #####
60         # TODO: Update self.xs.
61         # Hint: Call self.transition_model().
62         # Hint: You may find np.random.multivariate_normal useful.
63
64         us = np.random.multivariate_normal(u, self.R, self.M)
65         self.xs = self.transition_model(us, dt)
66
67         ##### Code ends here #####
68
69     def transition_model(self, us, dt):
70         """
71         Propagates exact (nonlinear) state dynamics.
72
73         Inputs:
74         us: np.array[M,2] - zero-order hold control input for each particle.
75         dt: float - duration of discrete time step.
76
77         Output:
78         g: np.array[M,3] - result of belief mean for each particle
79             propagated according to the system dynamics with
80             control u for dt seconds.
81         """
82         raise NotImplementedError("transition_model must be overridden by a subclass of EKF")
83
84     def measurement_update(self, z_raw, Q_raw):
85         """
86         Updates belief state according to the given measurement.
87
88         Inputs:
89         z_raw: np.array[I,2] - matrix of I rows containing (alpha, r)
90             for each line extracted from the scanner
91             data in the scanner frame.
92         Q_raw: [np.array[2,2]] - list of I covariance matrices corresponding

```

```

91         to each (alpha, r) row of z_raw.
92     Output:
93         None - internal belief state (self.x, self.ws) is updated in self.resample().
94     """
95     raise NotImplementedError("measurement_update must be overridden by a subclass of EKF")
96
97     def resample(self, xs, ws):
98         """
99         Resamples the particles according to the updated particle weights.
100
101         Inputs:
102             xs: np.array[M,3] - matrix of particle states.
103             ws: np.array[M,] - particle weights.
104
105         Output:
106             None - internal belief state (self.xs, self.ws) should be updated.
107         """
108         r = np.random.rand() / self.M
109
110         ##### Code starts here #####
111         # TODO: Update self.xs, self.ws.
112         # Note: Assign the weights in self.ws to the corresponding weights in ws
113         #       when resampling xs instead of resetting them to a uniform
114         #       distribution. This allows us to keep track of the most likely
115         #       particle and use it to visualize the robot's pose with self.x.
116         # Hint: To maximize speed, try to implement the resampling algorithm
117         #       without for loops. You may find np.linspace(), np.cumsum(), and
118         #       np.searchsorted() useful. This results in a ~10x speedup.
119
120         c = np.cumsum(ws)
121         m = np.arange(0, self.M)
122         u = c[-1] * (r + m / self.M)
123         idx = np.searchsorted(c, u, side='left') # indices corresponding to resampled states
124         self.xs = xs[idx]
125         self.ws = ws[idx]
126
127         ##### Code ends here #####
128
129     def measurement_model(self, z_raw, Q_raw):
130         """
131         Converts raw measurements into the relevant Gaussian form (e.g., a
132         dimensionality reduction).
133
134         Inputs:
135             z_raw: np.array[I,2] - I lines extracted from scanner data in
136                               rows representing (alpha, r) in the scanner frame.
137             Q_raw: [np.array[2,2]] - list of I covariance matrices corresponding
138                               to each (alpha, r) row of z_raw.
139
140         Outputs:
141             z: np.array[2I,] - joint measurement mean.
142             Q: np.array[2I,2I] - joint measurement covariance.
143         """
144         raise NotImplementedError("measurement_model must be overridden by a subclass of EKF")
145
146     class MonteCarloLocalization(ParticleFilter):
147
148     def __init__(self, x0, R, map_lines, tf_base_to_camera, g):
149         """
150         MonteCarloLocalization constructor.
151
152         Inputs:
153             x0: np.array[M,3] - initial particle states.
154             R: np.array[2,2] - control noise covariance (corresponding to dt = 1 second).
155             map_lines: np.array[J,2] - J map lines in rows representing (alpha, r).
156             tf_base_to_camera: np.array[3,] - (x, y, theta) transform from the
157                               robot base to camera frame.
158             g: float - validation gate.
159         """
160         self.map_lines = map_lines # Matrix of J map lines with (alpha, r) as rows
161         self.tf_base_to_camera = tf_base_to_camera # (x, y, theta) transform
162         self.g = g # Validation gate
163         super(self.__class__, self).__init__(x0, R)
164
165     def transition_model(self, us, dt):
166         """
167         Unicycle model dynamics.
168
169         Inputs:
170             us: np.array[M,2] - zero-order hold control input for each particle.
171             dt: float - duration of discrete time step.
172
173         Output:
174             g: np.array[M,3] - result of belief mean for each particle
175                               propagated according to the system dynamics with
176                               control u for dt seconds.
177         """
178
179         ##### Code starts here #####
180         # TODO: Compute g.
181         # Hint: We don't need Jacobians for particle filtering.

```

```

181     # Hint: A simple solution can be using a for loop for each particle
182     # and a call to tb.compute_dynamics
183     # Hint: To maximize speed, try to compute the dynamics without looping
184     # over the particles. If you do this, you should implement
185     # vectorized versions of the dynamics computations directly here
186     # (instead of modifying turtlebot_model). This results in a
187     # ~10x speedup.
188     # Hint: This faster/better solution does not use loop and does
189     # not call tb.compute_dynamics. You need to compute the idxs
190     # where abs(om) > EPSILON_OMEGA and the other idxs, then do separate
191     # updates for them
192
193     # indices corresponding to abs(om) <= EPSILON_OMEGA
194     idx1 = np.where(np.abs(us[:, 1]) <= EPSILON_OMEGA)[0]
195     x_new1 = self.xs[idx1, 0] + np.multiply(us[idx1, 0], np.cos(self.xs[idx1, 2])) * dt
196     y_new1 = self.xs[idx1, 1] + np.multiply(us[idx1, 0], np.sin(self.xs[idx1, 2])) * dt
197     theta_new1 = self.xs[idx1, 2] + us[idx1, 1] * dt
198     g1 = np.vstack((x_new1, y_new1, theta_new1)).T
199
200     # indices corresponding to abs(om) > EPSILON_OMEGA
201     idx2 = np.where(np.abs(us[:, 1]) > EPSILON_OMEGA)[0]
202     x_new2 = self.xs[idx2, 0] + np.multiply(np.divide(us[idx2, 0], us[idx2, 1]),
203                                             np.sin(self.xs[idx2, 2] + us[idx2, 1] * dt) -
204                                             np.sin(self.xs[idx2, 2]))
205     y_new2 = self.xs[idx2, 1] - np.multiply(np.divide(us[idx2, 0], us[idx2, 1]),
206                                             np.cos(self.xs[idx2, 2] + us[idx2, 1] * dt) -
207                                             np.cos(self.xs[idx2, 2]))
208     theta_new2 = self.xs[idx2, 2] + us[idx2, 1] * dt
209     g2 = np.vstack((x_new2, y_new2, theta_new2)).T
210
211     # final g matrix
212     g = np.zeros((self.M, 3))
213     g[idx1, :] = g1
214     g[idx2, :] = g2
215
216     ##### Code ends here #####
217
218     return g
219
220 def measurement_update(self, z_raw, Q_raw):
221     """
222     Updates belief state according to the given measurement.
223
224     Inputs:
225         z_raw: np.array[I,2] - matrix of I rows containing (alpha, r)
226                        for each line extracted from the scanner
227                        data in the scanner frame.
228         Q_raw: [np.array[2,2]] - list of I covariance matrices corresponding
229                        to each (alpha, r) row of z_raw.
230
231     Output:
232         None - internal belief state (self.x, self.ws) is updated in self.resample().
233     """
234     xs = np.copy(self.xs)
235     ws = np.zeros_like(self.ws)
236
237     ##### Code starts here #####
238     # TODO: Compute new particles (xs, ws) with updated measurement weights.
239     # Hint: To maximize speed, implement this without looping over the
240     # particles. You may find scipy.stats.multivariate_normal.pdf()
241     # useful.
242     # Hint: You'll need to call self.measurement_model()
243
244     vs, Q = self.measurement_model(z_raw, Q_raw)
245     ws = scipy.stats.multivariate_normal.pdf(vs, mean=None, cov=Q)
246
247     ##### Code ends here #####
248
249     self.resample(xs, ws)
250
251 def measurement_model(self, z_raw, Q_raw):
252     """
253     Assemble one joint measurement and covariance from the individual values
254     corresponding to each matched line feature for each particle.
255
256     Inputs:
257         z_raw: np.array[I,2] - I lines extracted from scanner data in
258                        rows representing (alpha, r) in the scanner frame.
259         Q_raw: [np.array[2,2]] - list of I covariance matrices corresponding
260                        to each (alpha, r) row of z_raw.
261
262     Outputs:
263         z: np.array[M,2I] - joint measurement mean for M particles.
264         Q: np.array[2I,2I] - joint measurement covariance.
265     """
266     ##### Code starts here #####
267     # TODO: Compute Q.
268     # Hint: You might find scipy.linalg.block_diag() useful
269
270     Q = scipy.linalg.block_diag(*Q_raw)

```

```

271
272     ##### Code ends here #####
273
274     return vs, Q
275
276 def compute_innovations(self, z_raw, Q_raw):
277     """
278     Given lines extracted from the scanner data, tries to associate each one
279     to the closest map entry measured by Mahalanobis distance.
280
281     Inputs:
282         z_raw: np.array[I,2] - I lines extracted from scanner data in
283                             rows representing (alpha, r) in the scanner frame.
284         Q_raw: np.array[I,2,2] - I covariance matrices corresponding
285                             to each (alpha, r) row of z_raw.
286
287     Outputs:
288         vs: np.array[M,2I] - M innovation vectors of size 2I
289                             (predicted map measurement - scanner measurement).
290     """
291     def angle_diff(a, b):
292         a = a % (2. * np.pi)
293         b = b % (2. * np.pi)
294         diff = a - b
295         if np.size(diff) == 1:
296             if np.abs(a - b) > np.pi:
297                 sign = 2. * (diff < 0.) - 1.
298                 diff += sign * 2. * np.pi
299             else:
300                 idx = np.abs(diff) > np.pi
301                 sign = 2. * (diff[idx] < 0.) - 1.
302                 diff[idx] += sign * 2. * np.pi
303         return diff
304
305     ##### Code starts here #####
306     # TODO: Compute vs (with shape [M x I x 2]).
307     # Hint: Simple solutions: Using for loop, for each particle, for each
308     #         observed line, find the most likely map entry (the entry with
309     #         least Mahalanobis distance).
310     # Hint: To maximize speed, try to eliminate all for loops, or at least
311     #         for loops over J. It is possible to solve multiple systems with
312     #         np.linalg.solve() and swap arbitrary axes with np.transpose().
313     #         Eliminating loops over J results in a ~10x speedup.
314     #         Eliminating loops over I results in a ~2x speedup.
315     #         Eliminating loops over M results in a ~5x speedup.
316     #         Overall, that's 100x!
317     # Hint: For the faster solution, you might find np.expand_dims(),
318     #         np.linalg.solve(), np.meshgrid() useful.
319
320     hs = self.compute_predicted_measurements()
321     I = z_raw.shape[0] # Number of observed lines
322     J = hs.shape[1] # Number of predicted lines
323
324     mat = np.zeros((self.M, I, J, 2)) # stores vij for each observation and each sample
325
326     # compute difference between observed and predicted values
327     alpha_diff = angle_diff(np.expand_dims(z_raw[:, 0], axis=(1, 2)),
328                             np.expand_dims(hs[:, :, 0], axis=0))
329     r_diff = np.expand_dims(z_raw[:, 1], axis=(1, 2)) - np.expand_dims(hs[:, :, 1], axis=0)
330     mat[:, :, :, 0] = alpha_diff.transpose(1, 0, 2)
331     mat[:, :, :, 1] = r_diff.transpose(1, 0, 2)
332
333     # changes dimensions of Q_raw from I x 2 x 2 to M x I x 2 x 2
334     Q_big = np.repeat(Q_raw[np.newaxis, :, :, :], self.M, axis=0)
335
336     # 'maha_dist' stores the Mahalanobis distance on its diagonals for each
337     # observation and each sample. Its dimensions are M x I x J x J
338     maha_dist = np.matmul(mat, np.matmul(np.linalg.inv(Q_big), mat.transpose(0, 1, 3, 2)))
339     diag_elements = np.diagonal(maha_dist, axis1=2, axis2=3)
340     idx = np.argmin(diag_elements, axis=2) # indices corresponding to minimum distance
341
342     # creating arrays to index from 'mat'
343     i1 = np.repeat(np.array(range(self.M)), I)
344     i2 = np.repeat(np.array([range(I)]), self.M, axis=0).flatten()
345
346     # indexing from 'mat' to get 'vs'
347     vs = mat[i1, i2, idx.flatten(), :].reshape((self.M, I, 2))
348
349     ##### Code ends here #####
350
351     # Reshape [M x I x 2] array to [M x 2I]
352     return vs.reshape((self.M, -1)) # [M x 2I]
353
354 def compute_predicted_measurements(self):
355     """
356     Given a single map line in the world frame, outputs the line parameters
357     in the scanner frame so it can be associated with the lines extracted
358     from the scanner measurements.
359
360     Input:
361         None

```



```

361         Output:
362         hs: np.array[M,J,2] - J line parameters in the scanner (camera) frame for M particles.
363         """
364         ##### Code starts here #####
365         # TODO: Compute hs.
366         # Hint: We don't need Jacobians for particle filtering.
367         # Hint: Simple solutions: Using for loop, for each particle, for each
368         #         map line, transform to scanner frame using tb.transform_line_to_scanner_frame()
369         #         and tb.normalize_line_parameters()
370         # Hint: To maximize speed, try to compute the predicted measurements
371         #         without looping over the map lines. You can implement vectorized
372         #         versions of turtlebot_model functions directly here. This
373         #         results in a ~10x speedup.
374         # Hint: For the faster solution, it does not call tb.transform_line_to_scanner_frame()
375         #         or tb.normalize_line_parameters(), but reimplement these steps vectorized.
376
377         J = self.map_lines.shape[1]
378         hs = np.zeros((self.M, J, 2))
379
380         # pose of the camera in the world frame
381         x_cam = self.xs[:, 0] + self.tf_base_to_camera[0] * np.cos(self.xs[:, 2]) - \
382             self.tf_base_to_camera[1] * np.sin(self.xs[:, 2])
383         y_cam = self.xs[:, 1] + self.tf_base_to_camera[0] * np.sin(self.xs[:, 2]) + \
384             self.tf_base_to_camera[1] * np.cos(self.xs[:, 2])
385         th_cam = self.xs[:, 2] + self.tf_base_to_camera[2]
386
387         # line parameters in the world frame
388         alpha = self.map_lines[0, :]
389         r = self.map_lines[1, :]
390
391         # broadcast 1D arrays to M x J matrices
392         alpha_MJ = np.tile(alpha, (self.M, 1))
393         r_MJ = np.tile(r, (self.M, 1))
394         x_cam_MJ = np.tile(x_cam.reshape(self.M, 1), J)
395         y_cam_MJ = np.tile(y_cam.reshape(self.M, 1), J)
396         th_cam_MJ = np.tile(th_cam.reshape(self.M, 1), J)
397
398         # line parameters in the camera frame
399         alpha_in_cam = alpha_MJ - th_cam_MJ
400         r_in_cam = r_MJ - np.multiply(x_cam_MJ, np.cos(alpha_MJ)) - \
401             np.multiply(y_cam_MJ, np.sin(alpha_MJ))
402
403         # normalizing line parameters
404         i1, i2 = np.where(r_in_cam < 0)
405         r_in_cam[i1, i2] = -r_in_cam[i1, i2]
406         alpha_in_cam[i1, i2] = np.pi + alpha_in_cam[i1, i2]
407         alpha_in_cam = (alpha_in_cam + np.pi) % (2 * np.pi) - np.pi
408
409         # final hs of dimensions M x J x 2
410         hs[:, :, 0] = alpha_in_cam
411         hs[:, :, 1] = r_in_cam
412
413         ##### Code ends here #####
414
415         return hs
416

```